

Variable Annealing Length and Parallelism in Simulated Annealing

Vincent A. Cicirello, Ph.D.

Professor of Computer Science / Behavioral Neuroscience

cicirelv@stockton.edu

<http://www.cicirello.org/>

STOCKTON



STOCKTON
UNIVERSITY

www.stockton.edu

Introduction

- We introduce a restart schedule of run lengths for an adaptive simulated annealer.
 - Eliminates need to know/predict available time a priori.
- We extend the restart schedule to parallel implementation.
- The variable annealing length restart schedule leads to improved anytime behavior early in the run.
- Discuss experiments with an NP-Hard scheduling problem with sequence-dependent setups.

Background: Simulated Annealing (SA)

- **Annealing:** process of slowly cooling a heated metal.
 - Heating metal allows shaping, while cooling slowly minimizes internal stress / more stable final state.
- **SA:** stochastic search inspired by annealing process.
 - Temperature parameter controls acceptance of neighbors.
 - High temperature early in run = random search
 - Low temperature late in run = stochastic hill climb
 - Cool too quickly = converge too soon to local optima
 - Cool too slowly = excessive run length

Background: SA Annealing Schedules

- Most common annealing schedules:
 - Exponential cooling: $T_{i+1} = \alpha T_i$
 - Linear cooling: $T_{i+1} = T_i - \Delta T$
- Boyan's Modified Lam Annealing Schedule (Boyan 1998):
 - Temperature fluctuates up/down to track a theoretical “ideal” neighbor acceptance rate
 - Acceptance rate decreases exponentially from 1.0 (random search) to 0.44 during first 15% of run, and held at 0.44 for next 50% of run.
 - Declines exponentially for last 35% of run (stochastic hill climb).

Related Work: Restarts

- Restarting SA:
 - Many have shown one long run of SA usually outperforms multiple short independent runs.
 - Effective SA restarts likely involves dependent runs.
 - E.g., Sadeh et al ('97) models expected cost improvement, abandons less promising runs, reanneals other prior runs.
- Restarting other forms of search quite effective:
 - Restarting backtracking CSP search (e.g., Luby schedule)

Related Work: Parallel SA

- Three categories of parallel SA:
 1. Parallel neighbor evaluations
 - E.g., speculative moves (Ludwin & Betz 2011)
 2. Parallel multistart (with dependent runs)
 - E.g., regular intervals sharing best solution among parallel instances (Jha & Menon, 2014)
 3. Optimizing subproblems in parallel
 - E.g., graph partitioning (Rahimian et al 2015)

Approach: Variable Annealing Length

- **Rationale:** long run of SA typically outperforms multiple short runs, but difficult to accurately predict available time.
- **Variable Annealing Length (VAL):**
 - Multistart SA with increasing run lengths.
 - The length of restart r , in number of SA evaluations is:
$$\text{MaxEvals}(r) = 1000 * 2^r$$
 - The multistart SA follows the sequence of run lengths: $\{1000, 2000, 4000, 8000, \dots\}$

Approach: Parallel VAL, version 0

- **Parallel Variable Annealing Length, v0 (P-VAL-0):**

- Assume N parallel instances of SA: $\{SA_0, SA_1, \dots, SA_{N-1}\}$
- The length of restart r of instance SA_i is:

$$\text{MaxEvals}_i(r) = 1000 * 2^{i+r*N}$$

- When $N=1$, P-VAL-0 reduces to VAL.
- Example, when $N=3$:
 - Instance 0 follows run lengths: $\{1000, 8000, 64000, \dots\}$
 - Instance 1 follows run lengths: $\{2000, 16000, 128000, \dots\}$
 - Instance 2 follows run lengths: $\{4000, 32000, 256000, \dots\}$

Approach: P-VAL-0's flaw

- P-VAL-0 is flawed:
 - Assuming long run superior to multiple short runs, benefit of parallelization is from completing long runs earlier.
 - As # parallel instances goes to infinity, the longest run completed by P-VAL-0 finishes twice as early as VAL.
 - For $N=4$ parallel instances, the longest run completed by P-VAL-0 finishes 1.875 times as early as VAL.
 - For $N=8$ parallel instances, ... finishes 1.992 times as early.
 - We hit the limiting behavior with relatively few parallel instances.

Approach: Parallel VAL

- **Parallel Variable Annealing Length (P-VAL):**

- Assume N parallel instances of SA: $\{SA_0, SA_1, \dots, SA_{N-1}\}$
- The length of restart r of instance SA_i is:

$$\text{MaxEvals}_i(r) = 1000 * 2^{(i \bmod 4) + r * \min(N, 4)}$$

- When $N \leq 4$, P-VAL = P-VAL-0; but when $N > 4$:
 - Instances $\{0, 4, 8, \dots\}$ have run lengths: $\{1000, 16000, 256000, \dots\}$
 - Instances $\{1, 5, 9, \dots\}$ have run lengths: $\{2000, 32000, 512000, \dots\}$
 - Instances $\{2, 6, 10, \dots\}$ have run lengths: $\{4000, 64000, 1024000, \dots\}$
 - Instances $\{3, 7, 11, \dots\}$ have run lengths: $\{8000, 128000, 2048000, \dots\}$

Experiments


- **Problem:** Scheduling with sequence-dependent setups, minimizing weighted tardiness
 - Used common benchmark set
 - Best exact solver, dynamic programming, > 2 weeks CPU time solving hardest instances. (Tanaka & Araki, 2013)
 - Variety of algorithms applied to problem: neighborhood search (Liao et al, 2012), iterated local search (Xu et al 2014), ACO (Liao & Juan 2007), among others

Experiments

- **Platform:** Ubuntu 14.04 server, 2 Xeon L5520 quad-core (2.27GHz), 32GB; Java 8, Java HotSpot 64-bit Server VM
- Sequential (N=1) and parallel (N=4, N=8) experiments.
- For each algorithm, 10 runs on each of 120 instances, logging best solution at 1 second intervals over 60 s.
- Compare VAL, P-VAL-0, and P-VAL to the following:
 - Fixed annealing length (FAL-x) of x of total run, with restarts
 - E.g., FAL-1=one long run, FAL- $\frac{1}{2}$ =run length half total time
 - Parallel fixed annealing length (P-FAL-x)


Sequential Results

- VAL dominates fixed annealing length early.
- No significant difference, at end of run (last 12s), between VAL and FAL-1 ($p>0.35$)
- Not visually evident: Single long fixed length does outperform restarts of short fixed length at end.




P-VAL vs P-VAL-0

- P-VAL-0 with $N > 4$ parallel runs doesn't improve performance
 - e.g., P-VAL-0 with $N=8$ no better than $N=4$ (green/pink)
- P-VAL continues to see performance gains for $N > 4$ parallel runs.





Parallel Results

- Results with $N=8$
- Unlike sequential case, P-VAL does not approx. performance of fixed length restarts at run end.
- P-VAL dominates approx. 80% of P-FAL's first run
 - For 48s vs P-FAL-1
 - For 24s vs P-FAL- $\frac{1}{2}$
 - For 12s vs P-FAL- $\frac{1}{4}$
 - For 6s vs P-FAL- $\frac{1}{8}$



Reannealing vs Independent Runs

- All results so far are for independent runs.
- Also considered restarts that reanneal the current best of run solution
 - Including across parallel runs.
- Sequential results:
 - No significant difference for VAL with reannealing vs independent runs.
 - Same is true for FAL- $\frac{1}{8}$ with reannealing vs independent runs.
 - Same is true in parallel



Conclusions

- Proposed a multistart SA, with variable annealing length
 - Eliminates need to know/predict available time for run
 - Issue not limited to Modified Lam (e.g., common exponential schedule becomes stochastic hill climb too soon if α too low).
 - Short early runs quickly find “good” solution, and increasing run length approximates final performance of long SA runs.
- Proposed parallel implementation
- Long fixed length runs better at end of run, but variable length restarts exhibits stronger anytime performance.

Questions

STOCKTON



STOCKTON
UNIVERSITY

www.stockton.edu