Randomizing Dispatch Scheduling Policies

Vincent A. Cicirello and **Stephen F. Smith**
The Robotics Institute
Carnegie Mellon University
5000 Forbes Avenue
Pittsburgh, PA 15213
{cicirello, sfs}@cs.cmu.edu

Abstract

The factory is a complex dynamic environment and scheduling operations for such an environment is a challenging problem. In practice, dispatch scheduling policies are commonly employed, as they offer an efficient and robust solution. However, dispatch scheduling policies are generally myopic, and as such they are susceptible to sub-optimal decision-making. In this paper, we attempt to improve upon results of such dispatch policies by introducing non-determinism into the decision-making process, and instead using a given policy as a baseline for biasing stochastic decisions. We consider the problem of weighted tardiness scheduling with sequence-dependent setups with unknown arrival times in a dynamic environment, and show that randomization of state-of-the-art dispatch heuristics for this problem in this manner can improve performance. Furthermore, we find that the "easier" the problem, the less benefit there is from randomization; the "harder" the problem, the more benefit. Our method of randomization is based on a model of the way colonies of wasps self-organize social hierarchies in nature.

Introduction

In dynamic factory environments, dispatch scheduling policies provide a practical, robust basis for managing execution. Scheduling decisions are made in an online manner only when needed, based on the current state of the factory. Dispatch policies make use of information about jobs such as expected processing time, setup time, due date, priority, etc., and are typically designed to optimize a given performance objective. Their virtue is their simplicity and insensitivity to environmental dynamics and for these reasons they are commonly employed. At the same time, the localized and myopic nature by which decisions are made under such schemes make them inherently susceptible to sub-optimal decision-making. (see (Morton & Pentico 1993) for a thorough review of dispatch scheduling).

In this paper, we explore the idea of randomizing dispatch policies to improve their performance. We start from the perspective that a given dispatch policy provides a good decision-making baseline, but it is not infallible and will sometimes make suboptimal choices. Hence, instead of unconditionally following its decisions, we use the dispatch policy to bias a stochastic selection process. Randomizing heuristics to improve decision-making performance is not a
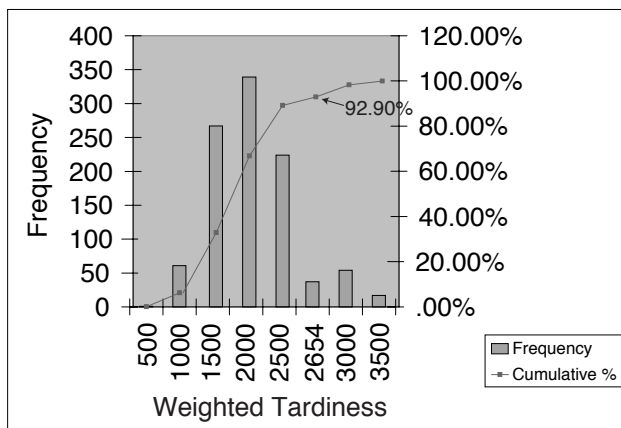

Figure 1: This histogram partitions 1000 trials of a randomized version of the dispatch policy COVERT on an instance of a weighted tardiness problem with sequence-dependent setups. Each bucket contains the trials with results in a given range of the weighted tardiness objective. COVERT produces a solution with weighted tardiness of 2654 for this instance. 92.9% of the random trials gave better solutions.

new idea; but to our knowledge this is the first attempt to do so in an online control setting.

Our approach hypothesizes that solution obtained using a dispatch policy is surrounded by a neighborhood of potentially "better" solutions. This hypothesis is based on a couple of observations. First, dispatch policies in practice tend to be quite sensitive to parameter settings and in fact they are frequently tuned to individual problem instances in experimental studies. (Such tuning is of course not possible in online settings.) Second, dispatch policies are often designed for idealized problems as they are adapted to accommodate more complex constraints this can result in more imprecise state estimation. Figure 1 shows the results of 1000 trials of a randomized dispatch policy designed to minimize weighted tardiness, partitioned according to observed weighted tardiness. Of these trials, 92.9% produced results better than the deterministic policy, giving empirical evidence of the potential of randomization.

We focus specifically on a weighted tardiness scheduling

problem with sequence dependent setups. In our problem, jobs arrive dynamically with specific due dates and priorities, and the goal is to minimize average weighted tardiness. The problem is further complicated by the fact that the machines used to process jobs, though they are multi-purpose, require some nontrivial amount of setup time to change-over from processing a job of one type to processing a job of another type.

Two of the best regarded dispatch policies for the weighted tardiness problem are COVERT (Vepsalainen & Morton 1987) and R&M (Rachamadugu & Morton 1982). Both are considered state-of-the-art (Morton & Pentico 1993); but neither were actually designed with sequence-dependent setups in mind. In fact, we are aware of only two reported dispatch heuristics for weighted tardiness scheduling under sequence-dependent setups (Raman, Rachamadugu, & Talbot 1989; Lee, Bhaskaran, & Pinedo 1997), both of which are derived by making (slightly different) modifications to the R&M dispatch policy to account for setup time.

In the remainder of this paper, we first review prior work in heuristic-biased stochastic scheduling methods. We then describe our approach of randomizing dispatch scheduling policies based on a model of wasp behavior. The paper then continues with a discussion of experimental results and conclusion.

Heuristic-Biased Search

Heuristic-biased random strategies have been used as effective iterative search strategies in a number of scheduling-related contexts. In (Watson *et al.* 1999), Watson et al. study the effects of problem structure on a number of scheduling algorithms for the flow-shop problem. They argue that more complex algorithms which work well on randomly generated benchmark problems often fail to live up to expectations when applied to real problems. The problem from their point of view is that real problems often contain structure of one sort or another that is not present in randomly generated benchmark problems. They hypothesize that more complex, carefully crafted algorithms can become over-tuned to the benchmarks. In their study, they compare a number of scheduling approaches for the flow-shop scheduling problem and show that as real-world problem-inspired structure is added to the problem, the faster and simpler stochastic algorithms are superior to the more complex "state-of-the-art" deterministic algorithms.

In (Bresina 1996), Bresina introduced a random restart search technique called heuristic-biased stochastic sampling (HBSS). HBSS performs a random search biased according to a heuristic for the problem. Bresina considered the scheduling of observations on a telescope. His approach began by ordering the observations according to a heuristic and giving each a rank according to the resulting order. The next observation would then be selected by a random process biased according to a function of these rankings. Different bias functions lead to more or less biasing in favor of the heuristic. So the choice of bias function depends on confidence in the heuristic.

Oddi and Smith (Oddi & Smith 1997) explored a related idea as the basis for solving a generalized job shop scheduling problem. One important distinction in this approach is acknowledgement of the fact that a heuristic may be more or less informed in different decision-making contexts, and hence the degree of confidence in the heuristic can vary from decision to decision. Rather than rely on a static bias function as is used in HBSS, Oddi and Smith define non-deterministic variants of search control heuristics that vary the degree of randomness as a function of how informed the heuristic is. A variant of this idea is also exploited by Cesta et al. (Cesta, Oddi, & Smith 2001) in solving a resource-constrained project scheduling problem.

In this paper, we take a similar approach but our goal is to make more effective control decisions rather than to broaden a search process. We bias our random decisions not by a function of a ranking given by the dispatch policy, but rather by a function of the value of the dispatch heuristic itself. In a control setting, it is particularly important to calibrate the degree of randomness of a decision to some measure of how discriminating the dispatch heuristic is in this decision context. Since we are making online decisions, we want nondeterminism when chances of finding a better solution than that proposed by the heuristic are best.

Scheduling Wasps

Our approach to randomization derives from a naturally-inspired computational model of the self-organization that takes place within a colony of wasps. (Theraulaz *et al.* 1991; Theraulaz, Bonabeau, & Deneubourg 1995; Bonabeau *et al.* 1997). In nature, a hierarchical social order among the wasps of the colony is formed through interactions among individual wasps of the colony. This emergent social order is a succession of wasps from the most dominant to the least dominant (analogous to a prioritization of jobs on a set of machines). In the model of Theraulaz et al., the results of these interactions are determined stochastically based on the "force" variables of the wasps involved. The probability of wasp 1 winning a dominance contest against wasp 2 is defined based on the force variables, F_1 and F_2, of the wasps as:

$$P\left(F_1, F_2\right) = \frac{F_1^2}{F_1^2 + F_2^2} \qquad (1)$$

This model can be directed mapped to the problem of prioritizing jobs in a queue, and as such provides a natural basis for the randomization of dispatch policies. In our "scheduling wasp" formulation, each job is represented by a wasp and the concept of a force variable is used to define job priority (i.e., the value that the dispatch policy in use assigns). The scheduling wasps then interact with each other to prioritize the jobs in the queue. This framework for dynamic scheduling was first introduced in (Cicirello & Smith 2001), where we considered the problem of sequencing jobs to maximize throughput under different and dynamically changing job mixes. Here we explore the use of this wasp model on due date problems, where dispatch-based solutions are more commonly employed.
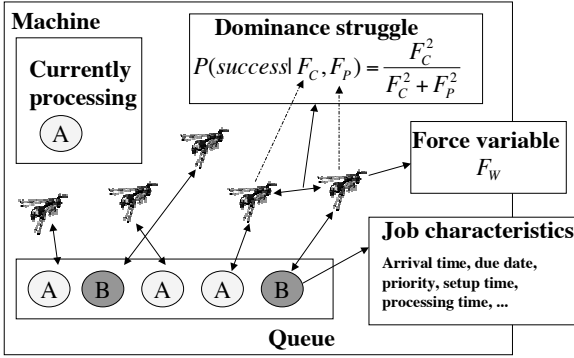
Scheduling Wasps

Figure 2: Scheduling wasps

To fully specify our scheduling wasp model for minimizing weighted tardiness with sequence dependent setups, we need to provide a definition of force. We will consider two alternatives: 1) COVERT (Vepsalainen & Morton 1987) and 2) R&M (Rachamadugu & Morton 1982). As mentioned earlier, neither of these dispatch policies was originally designed with sequence-dependent setups in mind. We have here modified both COVERT and R&M to account for setup time. The modification of R&M was made previously in (Raman, Rachamadugu, & Talbot 1989) and we have here made the equivalent modification to COVERT.

Noting this, the two alternative definitions of force are:

1. COVERT: $F_w = \frac{W_w}{T_w^p + T_w^s} \left(1.0 - \frac{(D_w - T_w^p - T_w^s - T_{now})^+}{h \, (T_w^p + T_w^s)}\right)+$

2. R&M: $F_w = \frac{W_w}{T_w^p + T_w^s} \exp\left(\frac{-(D_w - T_w^p - T_w^s - T_{now})^+}{h \, \overline{T^p}}\right)$

where T_w^p and T_w^s are the processing time and setup time of wasp w's job, D_w is the duedate, W_w is the weight, T_{now} is the current time, $\overline{T^p}$ is the average processing time, and $(A)^+ = \max\{A, 0\}$. The winner of a dominance contest in this context is determined stochastically in the same manner as in the model of real wasp behavior.

Dominance Tournaments

We have so far left open the question of what actually happens when a machine becomes available. In the typical dispatch scheduling approach, the job in the queue with the highest value of the dispatch heuristic is chosen next. Given the scheduling wasp formulation of the previous section, our system instead chooses the next job based on a tournament of dominance contests. We consider four dominance tournaments which form a progression from the most randomized to the least randomized. These four tournaments are:

1. Step-Ladder (W-V1): In this tournament, the scheduling wasps are seeded based on their current position in the queue. The last two wasps in the queue engage in a dominance contest. The winner then engages in a dominance contest with the next wasp and so forth along the length of the queue. As this occurs, the jobs associated with

the winning wasps move closer to the front of the queue. Whatever job is at the front of the queue when this process completes is chosen next by the machine.

2. Single-Elimination (W-V2): In this tournament, the scheduling wasps are seeded arbitrarily and take part in a single-elimination tournament. At each level of the tournament exactly half of the wasps are eliminated. This style tournament is biased slightly more in favor of whatever dispatch heuristic is used as force when compared to the previous. In W-V1, a job with a low force value at the front of the queue can potentially get lucky and win its only dominance contest; whereas in W-V2 winning a single dominance contest is not sufficient (unless there are only two jobs in the queue).

3. Double-Elimination (W-V3): This tournament is one step less random then the previous. Again the wasps are seeded arbitrarily. But this time, two simultaneous tournaments alternate rounds. That is, the primary tournament consists of undefeated wasps. The losers at each round move to the secondary tournament. In the secondary tournament, losers are eliminated. The winner of the primary tournament then gets two chances (has not lost yet) of eliminating the winner of the secondary tournament. If it is unsuccessful in both attempts then the winner of the secondary tournament is chosen by the machine. In this way, a wasp with a high force value that was unlucky and lost early in the tournament still has a chance of winning.

4. One-More-Chance (W-V4): This method begins by conducting a double-elimination tournament as in W-V3. If the winner also happens to have the highest value of the dispatch heuristic then it is chosen. Otherwise, it engages in a final dominance contest with the wasp with the highest value. And the winner of this dominance contest is taken next by the machine. So as you can see this is yet another step in the less random direction.

Degrees of Randomness

In describing the dominance tournaments above, we stated that the four models form a progression along a randomness scale beginning with the most random (W-V1) and ending with the least random (W-V4). Deterministic dispatch scheduling can be seen as a last step in this progression. To illustrate this progression, consider two examples. In Example 1 there are 4 jobs in the queue: A_1, A_2, A_3, A_4. Now consider that we have a dispatch heuristic H that assigns the following values to these jobs: $H(A_1) = H(A_2) = H(A_3) = 1$ and $H(A_4) = 3$. Let the Force F be defined by this heuristic. In Example 2, everything is the same except that $H(A_4) = 1.1$. In Table 1 we list the probability that job A_4 is chosen next by the machine in each example for each of the defined dispatching methods. In addition to showing the progressive increase in randomness through the wasp tournament variations, this Table illustrates the desired general property that the more certain the dispatch heuristic is, the less random the decision produced by our randomization methods; and the less certain the dispatch heuristic is the more random the decision.

Table 1: A couple examples of randomness progression.

	Example 1	Example 2
$P_{\text{roulette}}(A_4)$	0.7500	0.2874
$P_{\text{W-V1}}(A_4)$	0.7920	0.2939
$P_{\text{W-V2}}(A_4)$	0.810	0.2998
$P_{\text{W-V3}}(A_4)$	0.9331	0.3197
$P_{\text{W-V4}}(A_4)$	0.9933	0.6922
$P_{\text{dispatch}}(A_4)$	1.0	1.0

Experimental Design

All of the experiments that are presented here were performed in a simulated factory environment implemented in Java and executed on a Pentium III running Linux 5.2. We consider factories that produce two products (henceforth, Job Type A and Job Type B) as well as three products (Job Type A, Job Type B, and Job Type C). All machines in the factory are multi-purpose machines that can process any of these types (only single stage jobs are considered here). Experiments with one, two, and four machines are studied. In all cases, setup time to reconfigure a machine for a different job type is 30 time units.

In the two job type experiments, jobs are released to the factory floor dynamically according to three different product mixes. In each, arrival rates are defined by the probability a new job of each type is released during a given time unit. The arrival rates for the one machine problems are as follows:

- 50/50 mix: P(Job Type A) = 0.025, P(Job Type B) = 0.025
- 85/15 mix: P(Job Type A) = 0.04285, P(Job Type B) = 0.00715
- 100/0 mix: P(Job Type A) = 0.0665, P(Job Type B) = 0.0

Multiply these rates by 2 to get the arrival rates for the two machine problems and by 4 to get the rates for the four machine problems.

In the three job type experiments, the arrival rates are defined similarly for the one machine problem (make the same appropriate adjustments as above for two and four machine problems):

- 33/33/33 mix: P(Job Type A) = 0.0166, P(Job Type B) = 0.0166, P(Job Type C) = 0.0166
- 50/25/25 mix: P(Job Type A) = 0.025, P(Job Type B) = 0.0125, P(Job Type C) = 0.0125

These arrival rates correspond approximately to medium-to-heavily loaded factories.

When a new job is generated, its process time is 15 plus a Gaussian noise factor, its weight is drawn uniformly from the interval $[1, 20]$, and its duedate is drawn uniformly from one of the following intervals (where P is process time, W is weight, and T is current time):

- $[T, T + 4\,P]$ if $W > 16$
- $[T, T + 6\,P]$ if $12 < W \leq 16$
- $[T, T + 6.5\,P]$ if $8 < W \leq 12$
- $[T, T + 8\,P]$ if $W \leq 8$

Two sets of experiments were performed. The first using COVERT as the definition of the force variable for the scheduling wasps (the second using R&M). In each set of experiments, the four variations of the scheduling wasp selection method were compared to the base deterministic dispatch heuristic. For each combination of scheduling method and job mix, 100 simulations with different arrival sequences were performed and each simulation was at least 1000 time units in length (jobs stop arriving at time unit 1000 and simulation ends when all jobs have been processed).

Experimental Results

In Table 2 we see the average weighted tardiness of 100 simulations for the one, two, and four machine problems with two and three job types for various job mixes. The results in this table compare the variations of the scheduling wasps (using COVERT as force) with that of the dispatch heuristic COVERT. You will first notice that for the 100/0 job mix and for each of one, two, and four machines that COVERT performs best. The result is statistically significant according to an ANOVA with correlated samples. This is to be expected. The 100/0 mix is a single job type problem and therefore sequence-dependent setups are not an issue. This is what COVERT was especially designed for (i.e., no setups) so randomizing is almost certain to give us worse results. But do note that as you scan across the rows corresponding to the 100/0 job mix that the less random the scheduling wasp selection method, the closer the results are to that of COVERT. In fact, the result of a Tukey HSD test show in all cases no significance between W-V4 and COVERT.

Now turn your attention to the 85/15 job mix. This case is slightly less like a one job type problem and the sequence-dependent setups may play an active part. In the one, two, and four machine problems, the best results are again toward the COVERT end of the chart with W-V4 edging out COVERT slightly in all but the two machine problem. However, the one machine result is not statistically significant and in both the two and four machine case the pairwise result of W-V4 and COVERT is not significant according to the Tukey HSD test.

Next, examine the 50/50 job mix results. In this case, minimization of setups is much more crucial to performance. And in every case, we see that the most random of the wasp variations (WV-1) performs best and in all but the four machine problem the result is statistically significant. In the four machine problem, there are more degrees of freedom for the system to work with. That is, if a machine M1 is mislead by COVERT into taking a particular job J1 when J2 would have been better, there are three other machines M2, M3, and M4 that may finish whatever they are doing in time to pick up job J2 before it sits around long enough to cause any serious harm to the objective. In a sense the four machine problem is "easier" than the one and two machine problems and randomization seems to buy us less.

Based on the two job type problems using COVERT, it appears that the trend is that the more sequence-dependent setups are a factor to the problem, then the greater the benefit of randomization. And alternatively, the less sequence-

Table 2: Average weighted tardiness for different job mixes. COVERT is used as the Force definition for all variations of the scheduling wasp approach. The results are compared to COVERT. 95% confidence intervals and P-values from ANOVA tests with correlated samples are shown. HSD is the absolute difference between any two means required for significance at the 0.05 level according to the Tukey HSD test.

One Machine Problem							
	W-V1	W-V2	W-V3	W-V4	COVERT	P-value	HSD
50/50 mix	**2081.4±219.9**	2242.9±267.2	2489.6±273.1	2636.9±290.5	2637.8±288.9	<0.0001	267.6
85/15 mix	699.2±108.7	661.5±95.0	667.0±106.7	**650.7±99.8**	660.7±115.1	0.3915	–
100/0 mix	214.3±31.0	200.3±28.1	188.8±26.2	176.3±24.2	**172.5±23.8**	<0.0001	7.1
33/33/33 mix	**2834.5±217.1**	2970.8±221.0	3174.9±257.0	3215.6±259.5	3307.4±274.8	<0.0001	251.6
50/25/25 mix	**2582.4±226.1**	2765.4±249.8	2667.5±253.9	2865.8±279.3	2912.6±299.3	0.0009	243.1
Two Machine Problem							
	W-V1	W-V2	W-V3	W-V4	COVERT	P-value	HSD
50/50 mix	**1753.1±147.3**	1986.2±209.1	2089.1±224.4	2138.3±235.0	2129.6±247.1	0.0001	287.3
85/15 mix	437.0±56.1	438.0±54.4	409.9±49.8	398.7±49.0	**382.6±50.4**	0.0013	44.1
100/0 mix	137.0±20.9	127.6±17.6	120.5±16.5	114.9±15.7	**113.3±15.2**	<0.0001	5.3
33/33/33 mix	**2454.6±161.5**	2666.6±168.7	2804.4±206.6	2808.4±214.1	2808.3±202.3	<0.0001	237.8
50/25/25 mix	2399.9±156.4	**2391.5±176.9**	2456.3±204.2	2670.0±218.0	2529.8±215.6	0.0024	218.2
Four Machine Problem							
	W-V1	W-V2	W-V3	W-V4	COVERT	P-value	HSD
50/50 mix	**1490.5±118.3**	1558.2±146.2	1682.0±170.8	1595.4±190.6	1540.0±186.4	0.3466	–
85/15 mix	343.7±38.0	316.5±35.5	307.8±35.7	**291.5±34.2**	303.6±37.9	0.0005	33.8
100/0 mix	81.6±10.4	76.7±8.9	73.8±8.4	72.1±8.1	**70.8±8.0**	<0.0001	2.3
33/33/33 mix	2512.0±140.4	**2468.4±152.7**	2529.3±173.2	2665.2±175.0	2703.1±194.1	0.0137	223.2
50/25/25 mix	2122.9±116.7	**2093.0±140.7**	2176.7±173.8	2388.6±173.4	2315.7±180.9	0.0001	205.6

dependent setups are an issue the less randomized the heuristic should be. Also the more difficult the problem then the more randomization is needed.

We now turn to the results of the three job type problems (also appearing in Table 2) to see if this hypothesis holds. In both the 33/33/33 and 50/25/25 job mix problems, sequence-dependence should be more of an issue than the two job type problems and thus the problems should be more difficult. All results for both of the three job type mixes are statistically significant according to an ANOVA with correlated samples. Also for all numbers of machines and for all three job type mixes, the best result is either W-V1 or W-V2 (the more random end of the chart) with no significance between the W-V1 and W-V2 results according to Tukey HSD tests.

In Table 3 we see the complimentary results using R&M as the force value for the wasps as compared directly to using the dispatch heuristic R&M. In all cases there is no significant difference in the performance of R&M as compared to COVERT (from the previously discussed results) and there is also no significant difference between the best of the scheduling wasp variations using either dispatch heuristic as force. The overall trends seem to show the same thing as well. The more difficult the problem and the more that sequence-dependent setups are an issue, the more to be gained by randomization. The "easier" the problem, the better off is non-randomized dispatch scheduling.

Weighted Roulette Wheel

We have advocated our wasp model as a basis for randomization of dispatch policies. It is natural to ask whether or not there is a simpler way to randomize dispatch scheduling policies. One possibility that we consider is borrowed from the evolutionary computation community. The genetic algorithm (GA) selection strategy known as fitness proportional selection allocates to each individual of the population a chunk of a roulette wheel in size proportional to its fitness relative to the rest of the population. In a GA, this weighted roulette wheel would be spun n times to choose the n members of the successive generation.

As an alternative to our wasp formulations, we use a weighted roulette wheel. Each job j in the queue is allocated a portion of the wheel proportional to the square of its force value (i.e., F_j^2). And when a machine becomes available this wheel is spun once to select the next job from the queue. You can think of this as a generalization of a dominance contest between two wasps to that of a sort of "free-for-all" among all of the wasps in the nest. In this case, the probability that job j is processed by the available machine is given by:

$$P(F_j) = \frac{F_j^2}{\sum_{k=1}^{n} F_k^2} \qquad (2)$$

In Table 4, we see a comparison among W-V1 (using COVERT as Force), COVERT, and the weighted roulette wheel approach (Roulette) just described. We have chosen W-V1 for this comparison because it is the best performer overall from among the wasp variations. We can use this as

Table 3: Average weighted tardiness for different job mixes. R&M is used as the Force definition for all variations of the scheduling wasp approach. The results are compared to R&M. 95% confidence intervals and P-values from ANOVA tests with correlated samples are shown. HSD is the absolute difference between any two means required for significance at the 0.05 level according to the Tukey HSD test.

One Machine Problem							
	W-V1	W-V2	W-V3	W-V4	R&M	P-value	HSD
50/50 mix	**2133.5±234.2**	2307.1±267.4	2372.3±250.0	2737.8±298.1	2636.6±295.4	<0.0001	258.8
85/15 mix	691.6±110.6	678.2±98.6	666.6±100.9	**656.9±104.8**	665.5±116.0	0.7284	–
100/0 mix	215.8±31.2	204.5±28.6	190.5±26.0	178.9±24.5	**175.4±24.1**	<0.0001	7.2
33/33/33 mix	**2835.1±202.6**	3013.9±235.3	3132.7±251.6	3220.6±264.9	3325.2±272.4	<0.0001	260.1
50/25/25 mix	**2581.0±223.0**	2785.8±251.6	2711.0±256.3	2887.3±281.4	2917.2±296.1	0.0011	245.9
Two Machine Problem							
	W-V1	W-V2	W-V3	W-V4	R&M	P-value	HSD
50/50 mix	**1660.0±151.5**	1979.3±209.1	2083.5±206.6	2118.0±230.3	2111.0±240.1	<0.0001	281.7
85/15 mix	446.1±55.4	429.5±50.4	419.4±53.6	393.8±53.2	**386.4±50.0**	0.0009	44.3
100/0 mix	139.7±21.1	130.8±17.8	122.7±16.4	116.2±15.5	**115.3±15.3**	<0.0001	5.4
33/33/33 mix	**2554.5±168.7**	2611.0±155.0	2674.6±209.2	2890.7±211.2	2760.8±199.5	0.0014	240.9
50/25/25 mix	**2375.7±168.7**	2409.7±183.0	2449.1±193.1	2645.0±215.1	2563.9±222.4	0.0035	219.1
Four Machine Problem							
	W-V1	W-V2	W-V3	W-V4	R&M	P-value	HSD
50/50 mix	**1513.5±125.2**	1518.0±159.6	1677.1±195.4	1559.7±173.5	1540.9±191.2	0.4074	–
85/15 mix	344.2±40.2	309.1±33.5	305.9±34.4	**296.8±33.9**	300.7±37.8	0.0010	34.0
100/0 mix	83.5±10.4	79.2±8.9	75.6±8.4	73.8±8.1	**73.3±8.1**	<0.0001	2.2
33/33/33 mix	**2457.1±134.2**	2534.1±161.0	2598.1±182.6	2684.9±191.0	2637.1±197.4	0.0415	218.6
50/25/25 mix	2086.1±125.6	**2076.3±141.0**	2185.3±165.2	2210.0±175.5	2170.5±164.3	0.2299	–

Table 4: Average weighted tardiness for different job mixes. COVERT is used as the Force definition for the W-V1 approach. The results are compared to COVERT and the simpler randomizing scheme (Roulette). 95% confidence intervals and P-values from ANOVA tests with correlated samples are shown. HSD is the absolute difference between any two means required for significance at the 0.05 level according to the Tukey HSD test.

One Machine Problem				
	W-V1	Roulette	COVERT	P-value	HSD
50/50 mix	**2081.4±219.9**	2192.2±236.7	2637.8±288.9	<0.0001	256.0
85/15 mix	699.2±108.7	704.8±103.7	**660.7±115.1**	0.1713	–
100/0 mix	214.3±31.0	224.8±33.6	**172.5±23.8**	<0.0001	9.9
33/33/33 mix	**2834.5±217.1**	3000.4±230.8	3307.4±274.8	<0.0001	214.5
50/25/25 mix	**2582.4±226.1**	2741.8±253.2	2912.6±299.3	0.0021	219.2
Two Machine Problem				
	W-V1	Roulette	COVERT	P-value	HSD
50/50 mix	1753.1±147.3	**1735.6±159.5**	2129.6±247.1	<0.0001	219.3
85/15 mix	437.0±56.1	453.3±56.1	**382.6±50.4**	<0.0001	39.8
100/0 mix	137.0±20.9	144.0±22.0	**113.3±15.2**	<0.0001	6.9
33/33/33 mix	**2454.6±161.5**	2717.3±171.9	2808.3±202.3	<0.0001	191.5
50/25/25 mix	2399.9±156.4	**2360.7±160.4**	2529.8±215.6	0.0806	–
Four Machine Problem				
	W-V1	Roulette	COVERT	P-value	HSD
50/50 mix	**1490.5±118.3**	1604.4±127.7	1540.0±186.4	0.4463	–
85/15 mix	343.7±38.0	339.1±35.9	**303.6±37.9**	0.0025	29.6
100/0 mix	81.6±10.4	86.2±11.1	**70.8±8.0**	<0.0001	3.1
33/33/33 mix	2512.0±140.4	**2506.8±122.0**	2703.1±194.1	0.0127	177.1
50/25/25 mix	**2122.9±116.7**	2170.7±114.2	2315.7±180.9	0.0225	170.5

Table 5: Average weighted tardiness for different job mixes. R&M is used as the Force definition for the W-V1 approach. The results are compared to R&M and the simpler randomizing scheme (Roulette). 95% confidence intervals and P-values from ANOVA tests with correlated samples are shown. HSD is the absolute difference between any two means required for significance at the 0.05 level according to the Tukey HSD test.

One Machine Problem					
	W-V1	Roulette	R&M	P-value	HSD
50/50 mix	2133.5±234.2	**2128.8±236.6**	2636.6±295.4	<0.0001	257.9
85/15 mix	691.6±110.6	689.5±100.3	**665.5±116.0**	0.6257	–
100/0 mix	215.8±31.2	226.7±32.7	**175.4±24.1**	<0.0001	9.3
33/33/33 mix	**2835.1±202.6**	3044.6±225.3	3325.2±272.4	<0.0001	217.9
50/25/25 mix	**2581.0±223.0**	2757.7±255.4	2917.2±296.1	0.0012	213.5
Two Machine Problem					
	W-V1	Roulette	R&M	P-value	HSD
50/50 mix	**1660.0±151.5**	1804.2±157.4	2111.0±240.1	<0.0001	238.9
85/15 mix	446.1±55.4	474.4±62.5	**386.4±50.0**	<0.0001	40.0
100/0 mix	139.7±21.1	148.4±22.5	**115.3±15.3**	<0.0001	7.1
33/33/33 mix	**2554.5±168.7**	2741.6±171.0	2760.8±199.5	0.0222	193.4
50/25/25 mix	2375.7±168.7	**2344.1±153.3**	2563.9±222.4	0.0126	187.8
Four Machine Problem					
	W-V1	Roulette	R&M	P-value	HSD
50/50 mix	**1513.5±125.2**	1597.3±123.8	1540.9±191.2	0.6257	–
85/15 mix	344.2±40.2	318.5±36.6	**300.7±37.8**	0.0039	30.6
100/0 mix	83.5±10.4	88.8±11.1	**73.3±8.1**	<0.0001	3.06
33/33/33 mix	**2457.1±134.2**	2482.6±133.9	2637.1±197.4	0.0446	183.3
50/25/25 mix	**2086.1±125.6**	2199.3±132.2	2170.5±164.3	0.1680	–

a point of comparison to determine if the more complicated randomization scheme is doing anything purposeful. First note that for the one machine problem, W-V1 always outperforms Roulette. In the 100/0 job mix case, the difference is statistically significant based on a Tukey HSD test. Also in the 50/25/25 job mix, there is no pairwise significance via the Tukey HSD test between W-V1 and Roulette; however, there is significance between W-V1 and COVERT, but not between Roulette and COVERT. In all other mixes for the one machine problem, no pairwise statistical significance is seen between W-V1 and Roulette. In contrast, other than the 85/15 mix there is always pairwise statistical significance between COVERT and the others.

In the two machine problem W-V1 outperforms Roulette with statistical significance in both the 100/0 mix and 33/33/33 mix. In the other three mixes, no pairwise statistical significance was seen between W-V1 and Roulette. In the four machine problem, W-V1 again outperforms Roulette with statistical significance in the 100/0 mix and the 50/25/25 mix. The pairwise results in the other three job mixes showed no statistical significance between W-V1 and Roulette via Tukey HSD.

Overall, Roulette never does better than W-V1 with statistical significance and even in a couple cases when it does better without statistical significance COVERT outperforms both. W-V1, however, does outperform Roulette in a number of cases with statistical significance and in many of those cases also outperforms COVERT. The equivalent results using R&M show essentially the same thing as can be seen in Table 5. There is benefit to the "scheduling wasp model" approach to randomization.

Conclusion

In this paper we have shown that randomization of heuristics can be beneficial even in dynamic online situations. For the difficult problem of weighted tardiness scheduling with sequence-dependent setups, we have improved upon the performance of state-of-the-art weighted-tardiness dispatch policies by making decisions stochastically biased by the evaluations of these heuristics rather than using the dispatch policy directly. Our results show that the "easier" the problem and the less that sequence-dependent setups come into play, the less benefit is obtained through randomization. And furthermore the more that sequence-dependence is an issue and the more difficult the problem then the more benefit there is to randomization.

In showing this, we have developed and evaluated a progression of randomization methods for dispatch heuristics on a more/less random scale. These methods have at their foundation a model of self-organized wasp social hierarchies. Given their differential performance across our experimental design, one area of future research interest is to attempt to map characteristics of the current state of the problem to a degree of randomness required in the decision. This mapping can then allow us to incorporate a mix of the described approaches to make some decisions in a more or less stochastic manner. For instance, at some points in our simulation all jobs may be of a single type in which case it might

be perhaps better to use the dispatch policy directly; while at other points the mix might be closer to 50/50 (or some equivalent for more than two types) in which case a more random decision but biased by the heuristic might be better. In this manner, we may be able to combine the strengths of the deterministic and stochastic methods.

Another area we wish to explore is the use of our randomization scheme as an iterative search strategy, in much the same way as is done in (Bresina 1996; Oddi & Smith 1997) We are also exploring the application of our approach to other scheduling problems in both online and offline environments.

Acknowledgements

References

Bonabeau, E.; Sobkowski, A.; Theraulaz, G.; and Deneubourg, J. L. 1997. Adaptive task allocation inspired by a model of division of labor in social insects. In Lundh, D., and Olsson, B., eds., *Bio Computation and Emergent Computing*. World Scientific. 36–45.

Bresina, J. L. 1996. Heuristic-biased stochastic sampling. In *Proceedings of the Thirteenth National Conference on Artificial Intelligence and the Eighth Innovative Applications of Artificial Intelligence Conference, Volume One*, 271–278. AAAI Press.

Cesta, A.; Oddi, A.; and Smith, S. 2001. A constraint-based method for project scheduling with time windows. *Journal of Heuristics* To appear.

Cicirello, V. A., and Smith, S. F. 2001. Wasp nests for self-configurable factories. In *Agents 2001, Proceedings of the Fifth International Conference on Autonomous Agents*. ACM Press.

Lee, Y. H.; Bhaskaran, K.; and Pinedo, M. 1997. A heuristic to minimize the total weighted tardiness with sequence-dependent setups. *IIE Transactions* 29:45–52.

Morton, T. E., and Pentico, D. W. 1993. *Heuristic Scheduling Systems: With Applications to Production Systems and Project Management*. John Wiley and Sons.

Oddi, A., and Smith, S. F. 1997. Stochastic procedures for generating feasible schedules. In *Proceedings of the Fourteenth National Conference on Artificial Intelligence and Ninth Innovative Applications of Artificial Intelligence Conference*, 308–314. AAAI Press.

Rachamadugu, R. V., and Morton, T. E. 1982. Myopic heuristics for the single machine weighted tardiness problem. Working Paper 30-82-83, GSIA, Carnegie Mellon University, Pittsburgh, PA.

Raman, N.; Rachamadugu, R. V.; and Talbot, F. B. 1989. Real time scheduling of an automated manufacturing center. *European Journal of Operational Research* 40:222–242.

Theraulaz, G.; Goss, S.; Gervet, J.; and Deneubourg, J. L. 1991. Task differentiation in polistes wasp colonies: A model for self-organizing groups of robots. In *From Animals to Animats: Proceedings of the First International Conference on Simulation of Adaptive Behavior*, 346–355. MIT Press.

Theraulaz, G.; Bonabeau, E.; and Deneubourg, J. L. 1995. Self-organization of hierarchies in animal societies: The case of the primitively eusocial wasp polistes dominulus christ. *Journal of Theoretical Biology* 174:313–323.

Vepsalainen, A., and Morton, T. E. 1987. Priority rules and leadtime estimation for job shop scheduling with weighted tardiness costs. *Management Science* 33:1036–1047.

Watson, J. P.; Barbulescu, L.; Howe, A. E.; and Whitley, L. D. 1999. Algorithm performance and problem structure for flow-shop scheduling. In *Proceedings, Sixteenth National Conference on Artificial Intelligence (AAAI-99), Eleventh Innovative Applications of Artificial Intelligence Conference (IAAI-99)*, 688–695. AAAI Press.